

```

/*
 * Initialize shared buffer pool
 *
 * This is called once during shared-memory initialization (either in
the
 * postmaster, or in a standalone backend).
 */

```

- 设置4个bool变量: `bool foundBufs, foundDescs, foundIOLocks, foundBufCkpt;`
    - 将descriptor 对齐到cacheline boundary
    - 将lwlocks对齐到 cacheline boundary
    - 检查是否全部设置完毕, 如果没有, 初始化所有buffer的headers
    - 然后将所有的buffer链接到一起, 标记为unused
  - References [LWLockTranche::array\\_base](#), [LWLockTranche::array\\_stride](#), [Assert](#), [backend\\_flush\\_after](#), [buf](#), [BufferDesc::buf\\_id](#), [BufferBlocks](#), [BufferDescriptorGetContentLock](#), [BufferDescriptorGetIOLock](#), [BufferIO LWLockArray](#), [CLEAR\\_BUFFERTAG](#), [BufferDesc::freeNext](#), [FREEXT\\_END\\_OF\\_LIST](#), [GetBufferDescriptor](#), [i](#), [LWLockInitialize\(\)](#), [LWLockRegisterTranche\(\)](#), [LWTRANCHE\\_BUFFER\\_CONTENT](#), [LWTRANCHE\\_BUFFER\\_IO\\_IN\\_PROGRESS](#), [LWLockTranche::name](#), [NBuffers](#), [offsetof](#), [pg\\_atomic\\_init\\_u32\(\)](#), [ShmemInitStruct\(\)](#), [BufferDesc::state](#), [StrategyInitialize\(\)](#), [BufferDesc::tag](#), [BufferDesc::wait\\_backend\\_pid](#), and [WritebackContextInit\(\)](#).
  - Referenced by [CreateSharedMemoryAndSemaphores\(\)](#).
- Size BufferShmemSize(void)**
  - compute the size of shared memory for the buffer pool including **data pages, buffer descriptors, hash tables**, etc.
  - References [add\\_size\(\)](#), [mul\\_size\(\)](#), [NBuffers](#), [PG\\_CACHE\\_LINE\\_SIZE](#), and [StrategyShmemSize\(\)](#).

- o Referenced by [CreateSharedMemoryAndSemaphores\(\)](#).

```
o
    /*
     * Multiply two Size values, checking for overflow
     */
    Size
    mul_size(Size s1, Size s2);

    int          NBuffers = 1000;
```

## Variables

**BufferDescPadded\*** BufferDescriptors //buffer描述符

BufferDescPadded:

#include buf\_internals.sh // Internal definitions for *buffer manager* and the *buffer replacement*.

```
//Concurrent access to buffer headers has proven to be more efficient if
they're **cache line aligned**.So we force the start of the
BufferDescriptors array to be on a cache line boundary and force the
elements to be cache line sized.
```

```
#define BUFFERDESC_PAD_TO_SIZE (sizeof_void_p == 8 ? 64 : 1)
typedef union BufferDescPadded
{
    //BufferDesc -- shared descriptor/state data for a single shared buffer.
    BufferDesc bufferdesc;
    char pad[BUFFERDESC_PAD_TO_SIZE];
} BufferDescPadded;
```

```
typedef struct BufferDesc
{
    BufferTag tag;          /* ID of page contained in buffer */
    int buf_id;            /* buffer's index number (from 0) */

    /* state of the tag, containing flags, refcount and usagecount */
    pg_atomic_uint32 state;

    int wait_backend_pid;  /* backend PID of pin-count waiter
    */
    int freeNext;          /* link in freelist chain */

    LWLock content_lock;   /* to lock access to buffer contents */
} BufferDesc;
```

char\* BufferBlocks //Buffer实际的存储区域，类型为char\*

**LWLockMinimallyPadded\*** BufferIOLWLockArray = **NULL**

```
/* LWLock, minimally padded */
typedef union LWLockMinimallyPadded
{
    LWLock      lock;
    char        pad[LWLOCK_MINIMAL_SIZE];
} LWLockMinimallyPadded;
```

**LWLockTranche** BufferContentLWLockTranche

**LWLockTranche** BufferIOLWLockTranche

```

/*
 * Prior to PostgreSQL 9.4, every lightweight lock in the system was
stored
 * in a single array. For convenience and for compatibility with past
 * releases, we still have a main array, but it's now also permissible to
 * store LWLocks elsewhere in the main shared memory segment or in a
dynamic
 * shared memory segment.

**Each array of lwlocks forms a separate "tranche"**.

 *
 * It's occasionally necessary to identify a particular LWLock "by name";
e.g.
 * because we wish to report the lock to dtrace. We could store a name or
 * other identifying information in the lock itself, but since it's common
 * to have many nearly-identical locks (e.g. one per buffer) this would
end
 * up wasting significant amounts of memory. Instead, each lwlock stores
a
 * tranche ID which tells us which array it's part of. Based on that, we
can
 * figure out where the lwlock lies within the array using the data
structure
 * shown below; the lock is then identified based on the tranche name and
 * computed array index. We need the array stride because the array might
not
 * be an array of lwlocks, but rather some larger data structure that
includes
 * one or more lwlocks per element.
 */
typedef struct LWLockTranche
{
    const char *name;
    void      *array_base;
    Size       array_stride;
} LWLockTranche;

```

**WriteBackContext** BackendWritebackContext //保存写回时的上下文环境

```

/* struct forward declared in bufmgr.h */
typedef struct WritebackContext
{
    /* pointer to the max number of writeback requests to coalesce(合并) */
    int      *max_pending;

    /* current number of pending writeback requests */
    int      nr_pending;

    /* pending requests */
    PendingWriteback pending_writebacks[WRITEBACK_MAX_PENDING_FLUSHES];
} WritebackContext;

```

**CkptSortItem\*** CkptBufferIds //用于在checkpoint检查点时对每个file, sort它的buffer

```

/*
 * Structure to sort buffers per file on checkpoints.
 *
 * This structure is allocated **per buffer** in shared memory, so it
should be
 * kept as small as possible.
 */
typedef struct CkptSortItem
{
    Oid      tsId;
    Oid      relNode;
    ForkNumber forkNum;
    BlockNumber blockNum;
    int      buf_id;
} CkptSortItem;

```